

6.6 Lab 2: Ridge Regression and the Lasso

- 6.6.1 Ridge Regression
- 6.6.2 The Lasso

We will use the **glmnet** package in order to perform ridge regression and the lasso. The main function in this package is **glmnet()**, which can be used to fit ridge regression models, lasso models, and more. This function has a slightly different syntax from other model-fitting functions that we have encountered thus far in this book. In particular, we must pass in an **x** matrix as well as a **y** vector, and we do not use the **y ~ x** syntax. We will now perform ridge regression and the lasso in order to predict **Salary** on the **Hitters** data. Before proceeding ensure that the missing values have been removed from the data, as described in Section 6.5.

```
library(ISLR)
```

```
## Warning: package 'ISLR' was built under R version 3.4.4
```

```
Hitters = na.omit(Hitters)
x = model.matrix(Salary ~ ., Hitters[, -1])
y = Hitters$Salary
```

The **model.matrix()** function is particularly useful for creating **x**; not only does it produce a matrix corresponding to the 19 predictors but it also automatically transforms any qualitative variables into dummy variables. The latter property is important because **glmnet()** can only take numerical, quantitative inputs.

6.6.1 Ridge Regression

The **glmnet()** function has an **alpha** argument that determines what type of model is fit. If **alpha = 0** then a ridge regression model is fit, and if **alpha = 1**, then a lasso model is fit. We first fit a ridge regression model.

```
library(glmnet)
```

```
## Warning: package 'glmnet' was built under R version 3.4.4
```

```
## Loading required package: Matrix
```

```
## Loading required package: foreach
```

```
## Warning: package 'foreach' was built under R version 3.4.4
```

```
## Loaded glmnet 2.0-16
```

```
grid = 10^seq(10, -2, length = 100)
ridge.mod = glmnet(x, y, alpha = 0, lambda = grid)
```

By default the **glmnet()** function performs ridge regression for an automatically selected range of λ values. However, here we have chosen to implement the function over a grid of values ranging from $\lambda = 10^{10}$ to $\lambda = 10^{-2}$, essentially covering the full range of scenarios from the null model containing only the intercept, to the least squares fit. As we will see, we can also compute model fits for a particular value of λ that is not one of the original **grid** values. Note that by default, the **glmnet()** function standardizes the variables so that they are on the same scale. To turn off this default setting, use the argument **standardize = FALSE**.

Associated with each value of λ is a vector of ridge regression coefficients, stored in a matrix that can be accessed by **coef()**. In this case, it is a 20x100 matrix, with 20 rows (one for each predictors, plus an intercept) and 100 columns (one for each value of λ).

```
dim(coef(ridge.mod))
```

```
## [1] 20 100
```

We expect the coefficient estimates to be much smaller, in terms of l_2 norm, when a large value of λ is used, as compared to when a small value of λ is used. These are the coefficients when $\lambda = 11,498$, along with their l_2 norm:

```
ridge.mod$lambda[50]
```

```
## [1] 11497.57
```

```
coef(ridge.mod)[,50]
```

##	(Intercept)	(Intercept)	Hits	HmRun	Runs
##	420.533628801	0.000000000	0.142161648	0.535753108	0.237192308
##	RBI	Walks	Years	CAtBat	CHits
##	0.245400352	0.294764855	1.106768479	0.003145563	0.011707483
##	CHmRun	CRuns	CRBI	CWalks	LeagueN
##	0.087922960	0.023490502	0.024240472	0.025077345	0.031809908
##	DivisionW	PutOuts	Assists	Errors	NewLeagueN
##	-6.233613725	0.016680815	0.003066598	-0.011195017	0.272473462

```
sqrt(sum(coef(ridge.mod)[-1,50]^2))
```

```
## [1] 6.377993
```

In contrast, here are the coefficients when $\lambda = 705$, along with their l_2 norm. Note the much larger l_2 norm of the coefficients with this smaller value of λ .

```
ridge.mod$lambda[60]
```

```
## [1] 705.4802
```

```
coef(ridge.mod)[,60]
```

```
## (Intercept) (Intercept) Hits HmRun Runs
## 73.19638878 0.00000000 0.73352553 1.30217041 1.05007453
## RBI Walks Years CAtBat CHits
## 0.93206155 1.39062240 2.52015165 0.01096233 0.04728773
## CHmRun CRuns CRBI CWalks LeagueN
## 0.33786629 0.09449238 0.09826829 0.07092441 13.05069426
## DivisionW PutOuts Assists Errors NewLeagueN
## -54.60596556 0.12107727 0.02494856 -0.55777612 8.59917937
```

```
sqrt(sum(coef(ridge.mod)[-1,60]^2))
```

```
## [1] 56.91252
```

We can use the **predict()** function for a number of purposes. For instance, we can obtain the ridge regression coefficients for a new value of λ , say 50:

```
predict(ridge.mod, s = 50, type = 'coefficients')[1:20,]
```

```
## (Intercept) (Intercept) Hits HmRun Runs
## 23.18705218 0.00000000 1.41400571 -1.26017928 0.67493519
## RBI Walks Years CAtBat CHits
## 0.48417073 2.52681622 -5.87879262 0.00398633 0.10433894
## CHmRun CRuns CRBI CWalks LeagueN
## 0.61704523 0.21806578 0.22156158 -0.13369584 49.08322428
## DivisionW PutOuts Assists Errors NewLeagueN
## -120.86206156 0.24422330 0.08591212 -3.57430320 -12.43855470
```

We now split the samples into a training set and a test set in order to estimate the test error of ridge regression and the lasso. There are two common ways to randomly split a data set. The first is to produce a random vector of **TRUE**, **FALSE** elements and select the observations corresponding to **TRUE** for the training data. The second is to randomly choose a subset of numbers between 1 and n ; these can then be used as the indices for the training observations. The two approaches work equally well. We used the former method in Section 6.5.3. Here we demonstrate the latter approach.

We first set a random seed so that the results obtained will be reproducible.

```
set.seed(1)
train = sample(1:nrow(x), nrow(x)/2)
test = (-train)
y.test = y[test]
```

Next we fit a ridge regression model on the training set, and evaluate its MSE on the test set, using $\lambda = 4$. Note the use of the **predict()** function again. This time we get predictions for a test set, by replacing **type="coefficients"** with the **newx** argument.

```
ridge.mod = glmnet(x[train,], y[train], alpha = 0, lambda = grid, thresh = 1e-12)
ridge.pred = predict(ridge.mod, s = 4, newx = x[test,])
mean((ridge.pred - y.test)^2)
```

```
## [1] 107731.5
```

The test MSE is 101037. Note that if we had instead simply fit a model with just an intercept, we would have predicted each test observation using the mean of the training observations. In that case, we could compute the test set MSE like this:

```
mean((mean(y[train]) - y.test)^2)
```

```
## [1] 193253.1
```

We could also get the same result by fitting a ridge regression model with a very large value of λ . Note that **1e10** means 10^{10} .

```
ridge.pred = predict(ridge.mod, s = 1e10, newx = x[test,])
mean((ridge.pred - y.test)^2)
```

```
## [1] 193253.1
```

So fitting a ridge regression model with $\lambda = 4$ leads to a much lower test MSE than fitting a model with just an intercept. We now check whether there is any benefit to performing ridge regression with $\lambda = 4$ instead of just performing least squares regression. Recall that least squares is simply ridge regression with $\lambda = 0$.

In order for **glmnet()** to yield the exact least squares coefficients when $\lambda = 0$, we use the argument **exact=T** when calling the **predict()** function. Otherwise, the **predict()** function will interpolate over the grid of λ values used in fitting the **glmnet()** model, yielding approximate results. When we **exact=T**, there remains a slight discrepancy in the third decimal place between the output of **glmnet()** when $\lambda = 0$ and the output of **lm()**; this is due to numerical approximation on the part of **glmnet()**.

```
#ridge.pred = predict(ridge.mod, s = 0, newx = x[test,], exact = TRUE)
ridge.pred = predict(ridge.mod, s = 0, newx = x[test,])
mean((ridge.pred - y.test)^2)
```

```
## [1] 122927.1
```

```
lm(y ~ x, subset = train)
```

```
##
## Call:
## lm(formula = y ~ x, subset = train)
##
## Coefficients:
## (Intercept)  x(Intercept)      xHits      xHmRun      xRuns
##    191.2013         NA      1.2699     10.7712     -8.4450
##      xRBI      xWalks      xYears      xCAtBat      xCHits
##     1.4820     7.4691    -16.8688     -0.4000      0.8753
##    xCHmRun    xCRuns      xCRBI    xCWalks    xLeagueN
##    -0.9768     2.5616     -0.1611     -0.5112     71.8944
## xDivisionW    xPutOuts    xAssists    xErrors    xNewLeagueN
##   -113.9401      0.2859      0.1601     -0.8704     -13.1337
```


```
#predict(ridge.mod, s = 0, exact = TRUE, type = 'coefficients')[1:20,]
predict(ridge.mod, s = 0, type = 'coefficients')[1:20,]
```

```
## (Intercept) (Intercept)      Hits      HmRun      Runs
## 190.9617139  0.0000000  1.2793791  10.7956267  -8.4461360
##      RBI      Walks      Years      CAtBat      CHits
##  1.4634962  7.4756837 -16.9265952  -0.3966906   0.8624922
##    CHmRun    CRuns      CRBI    CWalks    LeagueN
## -0.9897737  2.5623244  -0.1541285  -0.5138003  71.6645361
##  DivisionW    PutOuts    Assists    Errors    NewLeagueN
## -114.1440109  0.2860706  0.1592226  -0.8777368  -13.0433027
```

In general, if we want to fit a (unpenalized) least squares model, then we should use the **lm()** function, since that function provides more useful outputs, such as standard errors and p-values for the coefficients.

In general, instead of arbitrarily choosing $\lambda = 4$, it would be better to use cross-validation to choose the tuning parameter λ . We can do this using the built-in cross-validation function, **cv.glmnet()**. By default, the function performs ten-fold cross-validation, though this can be changed using the argument **nfolds**. Note that we set a random seed first so our results will be reproducible, since the choice of the cross-validation folds is random.

```
set.seed(1)
cv.out = cv.glmnet(x[train,], y[train], alpha = 0)
plot(cv.out)
```



```
bestlam = cv.out$lambda.min
bestlam
```

```
## [1] 255.0435
```

Therefore, we see that the value of λ that results in the smallest cross-validation error is 212. What is the test MSE associated with this value of λ ?

```
ridge.pred = predict(ridge.mod, s = bestlam, newx = x[test,])
mean((ridge.pred - y.test)^2)
```

```
## [1] 96046.9
```

This represents a further improvement over the test MSE that we got using $\lambda = 4$. Finally, we refit our ridge regression model on the full data set, using the value of λ chosen by cross-validation, and examine the coefficient estimates.

```
out = glmnet(x, y, alpha = 0)
predict(out, type = 'coefficients', s = bestlam)[1:20,]
```


##	(Intercept)	(Intercept)	Hits	HmRun	Runs
##	17.10374672	0.00000000	0.99473300	0.39155947	1.16249574
##	RBI	Walks	Years	CAtBat	CHits
##	0.92499772	1.75941083	0.65724377	0.01136837	0.06184547
##	CHmRun	CRuns	CRBI	CWalks	LeagueN
##	0.43057438	0.12294381	0.13013453	0.03962605	24.60058036
##	DivisionW	PutOuts	Assists	Errors	NewLeagueN
##	-86.21289069	0.18208758	0.04168183	-1.54423302	8.41679514

As expected, none of the coefficients are zero – ridge regression does not perform variable selection!

6.6.2 The Lasso


We saw that ridge regression with a wise choice of λ can outperform least squares as well as the null model on the **Hitters** data set. We now ask whether the lasso can yield either a more accurate or a more interpretable model than ridge regression. In order to fit a lasso model, we once again use the **glmnet()** function; however, this time we use the argument **alpha = 1**. Other than that change, we proceed just as we did in fitting a ridge model.

```
lasso.mod = glmnet(x[train,], y[train], alpha = 1, lambda = grid)
plot(lasso.mod)
```



We can see from the coefficient plot that depending on the choice of tuning parameter, some of the coefficients will be exactly equal to zero. We now perform cross-validation and compute the associated test error.

```
set.seed(1)
cv.out = cv.glmnet(x[train,], y[train], alpha = 1)
plot(cv.out)
```



```
bestlam = cv.out$lambda.min
lasso.pred = predict(lasso.mod, s = bestlam, newx = x[test,])
mean((lasso.pred - y.test)^2)
```

```
## [1] 100743.4
```

This is substantially lower than the test set MSE of the null model and of least squares, and very similar to the test MSE of ridge regression with λ chosen by cross-validation.

However, the lasso has a substantial advantage over ridge regression in that the resulting coefficient estimates are sparse. Here we see that 12 of the 19 coefficient estimates are exactly zero. So the lasso model with λ chosen by cross-validation contains only seven variables.

```
out = glmnet(x, y, alpha = 1, lambda = grid)
lasso.coef = predict(out, type = 'coefficients', s = bestlam)[1:20,]
lasso.coef
```


##	(Intercept)	(Intercept)	Hits	HmRun	Runs
##	18.5394844	0.0000000	1.8735390	0.0000000	0.0000000
##	RBI	Walks	Years	CAtBat	CHits
##	0.0000000	2.2178444	0.0000000	0.0000000	0.0000000
##	CHmRun	CRuns	CRBI	CWalks	LeagueN
##	0.0000000	0.2071252	0.4130132	0.0000000	3.2666677
##	DivisionW	PutOuts	Assists	Errors	NewLeagueN
##	-103.4845458	0.2204284	0.0000000	0.0000000	0.0000000

```
lasso.coef[lasso.coef != 0]
```

##	(Intercept)	Hits	Walks	CRuns	CRBI
##	18.5394844	1.8735390	2.2178444	0.2071252	0.4130132
##	LeagueN	DivisionW	PutOuts		
##	3.2666677	-103.4845458	0.2204284		